

NAVIGATORS — MASTER STRICT FULL-STACK IMPLEMENTATION, REPAIR & VALIDATION PROMPT

You are the lead software architect, ML engineer, navigation engineer, backend engineer, and QA engineer responsible for taking the existing **Navigators / SIH26168** repository to a fully working, integrated, demonstrable system.

Do not treat this as a documentation task.

Do not merely explain what should be implemented.

You must inspect the existing repository, understand what already exists, identify defects and missing components, implement the required functionality, integrate all subsystems, run tests and benchmarks, fix failures, and leave the repository in a coherent working state.

The final goal is a **smartphone-only intelligent dead-reckoning navigation system**, with a robust **Python-based ML/navigation backend/core engine**, local/offline operation, and a usable frontend.

1. PROJECT OBJECTIVE

Project:

Navigators

Problem Statement:

SIH26168 — AI/ML Based Intelligent Dead Reckoning System for Seamless Navigation

The system must maintain useful vehicle navigation when GNSS becomes unavailable.

The system must use:

- smartphone accelerometer
- smartphone gyroscope
- smartphone GNSS when available
- AI/ML-based motion/velocity estimation
- inertial navigation
- 15-state EKF
- NHC
- ZUPT where appropriate
- offline map matching

The system must NOT require:

- OBD-II
- vehicle speedometer
- CAN bus
- external IMU

- FOG-IMU
- dedicated navigation hardware
- dedicated vehicle hardware

The smartphone itself is the sensing platform.

2. CORE ARCHITECTURE DECISION

Use **Python** as the primary engineering language for the backend/core navigation engine and ML pipeline.

Use Python for:

- sensor processing
- preprocessing
- calibration
- coordinate transformations
- alignment
- AI inference
- EKF
- NHC
- ZUPT
- GNSS state machine
- map matching
- replay
- simulation
- benchmarking
- testing
- local API/backend

Use PyTorch for model development/training.

Use FastAPI where a local API is useful.

Use WebSocket where continuous real-time sensor/state streaming is useful.

The backend **MUST** be designed as a **local software service**, not a mandatory cloud service.

3. LOCAL BACKEND VS CLOUD

This distinction is critical.

A Python backend is allowed.

A local Python backend is encouraged where it improves reliability and integration.


A remote/cloud backend is NOT required for navigation.

The final system should be capable of operating without:

- internet
- cloud inference

- remote APIs
- remote databases
- remote map services
- online routing
- online model downloads

The architecture may therefore be:



During development, the Python backend may run on the development computer.

For real-device deployment, investigate the most practical way to run the same core engine locally/on-device without redesigning the navigation algorithms.

Do NOT pretend that a laptop-hosted FastAPI server is “on-device.”

Clearly distinguish:

- development architecture
- demonstration architecture
- final deployment architecture

4. FIRST ACTION — COMPLETE REPOSITORY AUDIT

Before changing major code, inspect the entire repository.

Determine:

- project structure
- backend
- frontend
- Python modules
- ML modules
- model files
- datasets
- configuration
- APIs
- WebSocket implementation
- sensor ingestion
- simulator
- replay system

- EKF
- filtering
- alignment
- map matching
- tests
- benchmarks
- deployment scripts
- Docker configuration if present
- dependency management
- documentation
- obsolete hardware assumptions
- cloud assumptions

Do not trust documentation claims.

Do not trust progress reports.

Do not trust “completed” labels.

Verify actual implementation.

5. BASELINE BEFORE FIXING

Run the existing system before making destructive changes.

Run:

- unit tests
- integration tests
- backend
- frontend
- simulator
- replay
- ML inference
- available benchmarks

Record:

- errors
- failing tests
- missing dependencies
- broken imports
- crashes
- NaNs
- numerical instability
- latency
- memory issues
- incorrect trajectories
- coordinate-frame problems
- sensor-rate problems
- API problems
- WebSocket problems
- map matching problems

Do not delete failing tests.

6. FINAL SOFTWARE STACK

Prefer a clean stack similar to:

Backend

- Python
- FastAPI
- Pydantic
- NumPy
- SciPy
- PyTorch
- ONNX Runtime where useful
- WebSocket
- local storage/database only where justified

Navigation

- NumPy/SciPy
- custom tested EKF
- filtering
- coordinate transformations
- NHC
- ZUPT
- map matching

ML

- PyTorch
- TCN or LSTM
- reproducible training
- ONNX export where deployment requires it

Frontend

Use the existing frontend if viable.

If the existing frontend is incomplete, repair it rather than unnecessarily replacing it.

It must provide:

- map
- position
- navigation mode
- GNSS status
- velocity
- heading
- system health
- DR state

- relevant telemetry
-

7. REPOSITORY CLEANUP

Remove or isolate obsolete assumptions involving:

- external FOG-IMU
- external IMU
- OBD-II
- CAN
- dedicated vehicle hardware
- mandatory cloud backend

Do not blindly delete useful development code.

Separate it into optional/development modules if it remains useful.

The final system must clearly state:

No dedicated external hardware is required.

8. SENSOR INGESTION

Build a robust sensor interface.

Support:

- accelerometer
- gyroscope
- GNSS
- timestamps
- sensor permissions
- sensor availability
- dropped samples
- variable sampling rate

Never assume a sensor produces an exact frequency.

Measure actual frequency.

Normalize all sensor data into a common internal representation.

Every sample must carry a trustworthy timestamp.

9. SENSOR SYNCHRONIZATION

Implement timestamp-based synchronization.

Handle:

- asynchronous sensors
- different frequencies
- missing samples
- jitter
- duplicate samples
- timestamp discontinuities

Do not synchronize sensors merely by array index.

10. UNITS AND COORDINATE FRAMES

Explicitly define:

- phone frame
- vehicle frame
- navigation frame
- ENU/NED convention
- acceleration units
- gyro units
- velocity units
- position representation
- angular units

Every transformation must be explicit.

Create tests for:

- axis permutations
 - sign changes
 - rotations
 - known synthetic motion
-

11. PHONE-TO-VEHICLE ALIGNMENT

Implement robust phone-to-vehicle alignment.

The system should tolerate realistic phone mounting orientation.

Handle:

- roll
- pitch
- yaw

Test identical vehicle trajectories using multiple simulated phone orientations.

Navigation output should remain reasonably consistent.

12. PREPROCESSING

Implement:

Nonlinear median filter

for:

- shocks
- potholes
- impulse noise

Butterworth low-pass filter

for:

- high-frequency vibration
- sensor noise

Document:

- order
- cutoff
- sampling assumptions
- startup behavior
- state handling
- latency

Do not use future samples in live processing.

13. AI VELOCITY ESTIMATION

Implement the required lightweight temporal model.

Preferred:

TCN

Alternative:

LSTM

The model should estimate:

2D vehicle velocity

rather than merely scalar speed.

Define:

- input features

- window length
- normalization
- output frame
- output units
- inference frequency

Validate model outputs before feeding them into the EKF.

14. ML TRAINING PIPELINE

Create a reproducible training pipeline.

Include:

- dataset loading
- preprocessing
- trajectory-level splitting
- normalization
- augmentation
- training
- validation
- checkpointing
- evaluation

Prevent:

- drive leakage
- route leakage
- overlapping-window leakage
- normalization leakage

Do not split random windows from the same trajectory into train/test.

15. ML AUGMENTATION

Use realistic augmentation:

- noise
- bias
- drift
- orientation perturbation
- vibration
- temporal jitter
- missing samples where appropriate

Verify augmentation improves generalization.

16. PYTORCH / ONNX VALIDATION

If ONNX is used:

Run automated parity tests between:

PyTorch reference

and

ONNX Runtime

using identical input tensors.

Measure:

- maximum absolute difference
- mean absolute difference
- output shape agreement

Do not consider export successful merely because a .onnx file was generated.

17. 15-STATE EKF

Implement or repair the 15-state EKF.

Explicitly define:

- state vector
- transition model
- process model
- covariance
- process noise
- measurement models
- Jacobians
- initialization
- update equations

Validate numerical stability.

Detect:

- NaN
- Inf
- covariance explosion
- singular matrices
- invalid inversions
- invalid innovation

Use robust numerical methods.

18. EKF MEASUREMENT FUSION

The EKF must be capable of incorporating:

- GNSS
- AI-derived 2D velocity
- NHC
- ZUPT

depending on system state.

Do not blindly trust any measurement.

Implement reasonable gating and confidence handling.

19. GNSS STATE MACHINE

Implement:

```
GNSS_HEALTHY
  ↓
GNSS_DEGRADED
  ↓
GNSS_OUTAGE
  ↓
DEAD_RECKONING
  ↓
GNSS_REACQUISITION
  ↓
GNSS_HEALTHY
```

Use:

- temporal persistence
- hysteresis
- accuracy checks
- innovation gating

Do not switch state from one noisy GNSS sample.

20. GNSS OUTLIER REJECTION

Detect unrealistic GNSS measurements.

Consider:

- reported accuracy
- innovation
- velocity consistency
- temporal consistency
- physically plausible movement

Reject or downweight obvious outliers.

21. DEAD-RECKONING MODE

During GNSS outage:

Do NOT freeze the last location.

Do NOT draw a predetermined line.

Do NOT fake movement.

Use:

IMU
↓
Filtering
↓
Alignment
↓
AI 2D velocity
↓
EKF
↓
NHC/ZUPT
↓
Dead reckoning

The resulting trajectory must come from the estimator.

22. NHC

Implement appropriate non-holonomic constraints for ground vehicles.

Use them to constrain physically implausible lateral/vertical motion.

Do not over-constrain legitimate turning dynamics.

Test:

- straight driving
 - turns
 - acceleration
 - braking
 - stationary periods
-

23. ZUPT

Implement zero-velocity updates only when the vehicle is actually stationary.

Build a stationary detector.

Test:

- traffic light
- parking
- stopped vehicle
- low-speed movement
- moving vehicle

Prevent false ZUPT application.

24. OFFLINE MAP MATCHING

Implement map matching using locally available data.

No runtime calls to:

- Google Maps
- Mapbox
- OSRM
- Valhalla
- online OSM APIs
- remote routing APIs

Use:

- geometric matching
- HMM matching
- or another properly justified offline method

Map matching must correct/validate the estimated trajectory.

It must not manufacture a fake route.

25. LOCAL MAP STORAGE

Design a practical offline map architecture.

Possible structure:

```
offline_maps/  
  region/  
    roads/  
    nodes/  
    edges/  
    spatial_index/
```

Use an efficient spatial index.

Avoid loading unnecessarily large geographic datasets into memory.

26. FASTAPI

Use FastAPI as the local service layer where useful.

Possible endpoints:

```
GET /health
GET /status
POST /session/start
POST /session/stop
POST /sensor/batch
GET /navigation/state
GET /navigation/trajectory
GET /metrics
```

Use WebSocket for continuous streaming where appropriate.

Do not send every individual sensor value through expensive synchronous HTTP requests if WebSocket/batching is more appropriate.

27. BACKEND API DESIGN

Keep API models explicit.

Separate:

- raw sensor data
- processed sensor data
- AI prediction
- EKF state
- navigation state
- map-matching result
- diagnostics

Use Pydantic schemas or equivalent validation.

Reject malformed sensor packets.

28. REAL-TIME PROCESSING

Design the backend so that sensor processing does not block the API/UI.

Separate:

- ingestion
- processing
- inference
- navigation
- output streaming

Use appropriate:

- queues
- asynchronous processing

- worker threads/processes

where needed.

Measure end-to-end latency.

Do not claim real-time performance without measurement.

29. FRONTEND

The frontend must display the actual backend navigation state.

Do not implement a fake UI trajectory.

Display:

- map
- current position
- trajectory
- GNSS state
- DR state
- velocity
- heading
- confidence
- sensor status
- connection status

The frontend must clearly distinguish:

GNSS navigation

from

dead reckoning

30. LIVE SENSOR PIPELINE

Implement:

```
Smartphone
↓
Sensor interface
↓
FastAPI/WebSocket
↓
Python processing engine
↓
AI
↓
EKF
↓
NHC/ZUPT
↓
```

Map matching
↓
Navigation state
↓
WebSocket
↓
Frontend

For local demonstrations, all components should run on the local network or same machine/device as appropriate.

No cloud service should be required.

31. OFFLINE OPERATION

The system must have a clearly defined offline mode.

No internet should be required for:

- AI
- EKF
- filtering
- NHC
- ZUPT
- navigation
- map matching

If the frontend requires static assets, cache/package them locally.

If the Python backend is part of the selected deployment architecture, package it locally rather than hosting it remotely.

32. REPLAY SYSTEM

Build a deterministic replay engine.

Input:

recorded accelerometer
recorded gyroscope
recorded GNSS
timestamps
ground truth where available

Output:

estimated trajectory
velocity
heading
GNSS state
DR state
errors
latency

Allow artificial GNSS outages.

33. REQUIRED TEST SCENARIOS

Test:

Normal GNSS

GNSS + IMU + AI.

50 m outage

Target:

<5 m drift where validated by suitable ground truth.

1 km outage at approximately 60 km/h

Target:

<100 m drift where validated by suitable ground truth.

GNSS recovery

Verify smooth re-acquisition.

Phone orientation

Test different phone mounting orientations.

GNSS outliers

Inject incorrect GNSS measurements.

Sensor noise

Inject realistic IMU noise and bias.

34. ABLATION TESTS

Compare:

1. raw IMU
2. filtered IMU
3. AI velocity
4. AI + EKF
5. AI + EKF + NHC

6. AI + EKF + NHC + ZUPT

7. complete system + map matching

Measure actual improvement.

35. PERFORMANCE METRICS

Measure:

Accuracy

- final position error
- RMSE
- mean error
- maximum error
- drift
- drift percentage
- velocity error
- heading error

Runtime

- sensor rate
- AI inference rate
- AI latency
- EKF latency
- map matching latency
- total processing latency
- CPU
- RAM
- model size

Reliability

- dropped samples
 - invalid measurements
 - NaNs
 - Inf values
 - EKF failures
 - API failures
 - WebSocket disconnects
 - offline failures
-

36. NO FABRICATION

NEVER fabricate:

- accuracy

- drift
- latency
- sensor rate
- CPU usage
- memory
- benchmark results
- test results

If data is unavailable:

BLOCKED FOR VALIDATION

must be reported.

37. TEST SUITE

Implement:

Unit tests

For every mathematical component.

Integration tests

For every subsystem boundary.

End-to-end tests

For complete navigation.

Regression tests

For every discovered bug.

Run the entire test suite after major changes.

38. NO TEST CHEATING

Never:

- delete failing tests
- weaken assertions
- hardcode trajectories
- hardcode positions
- hardcode AI outputs
- force map snapping
- hide exceptions
- suppress NaNs
- replace real algorithms with mocks in production

39. LOGGING AND OBSERVABILITY

Implement structured logging.

At minimum expose:

- sensor rate
- GNSS state
- AI prediction
- EKF state
- covariance health
- innovation
- NHC status
- ZUPT status
- map matching status
- processing latency
- errors

Provide a debug mode for development.

Do not overwhelm production logs with raw high-frequency sensor data.

40. CONFIGURATION

Move tunable parameters into explicit configuration.

Examples:

- filter parameters
- model path
- model window
- EKF noise
- GNSS thresholds
- NHC thresholds
- ZUPT thresholds
- map-matching parameters
- logging level

Avoid scattered magic numbers.

41. DEPENDENCY MANAGEMENT

Create a reproducible environment.

Provide:

- requirements file or pyproject
- installation instructions

- model setup
- map-data setup
- test commands
- run commands

The repository should be reproducible on a clean development environment.

42. DO NOT REWRITE EVERYTHING UNNECESSARILY

Preserve working code.

Fix incrementally.

Before replacing a subsystem:

1. understand it
2. test it
3. identify the actual defect
4. determine whether repair is sufficient
5. replace only if justified

Do not perform a giant blind rewrite.

43. DEVELOPMENT VS FINAL DEPLOYMENT

Explicitly distinguish:

Development

Python backend may run on:

- laptop
- desktop
- local development machine

Demonstration

Use a local controlled environment.

Final deployment

Determine the most practical smartphone deployment strategy for the Python navigation engine and frontend.

If pure Python cannot directly run in the target browser/mobile environment, do NOT hide that limitation.

Evaluate appropriate options for packaging the Python core for the target platform.

Possible approaches may include:

- native mobile Python embedding
- Python mobile runtime
- compiled/native numerical components
- local application service
- hybrid architecture

Choose based on actual repository constraints and demonstrate the chosen path.

Do not falsely label a laptop-hosted backend as on-device processing.

44. ARCHITECTURE SHOULD REMAIN MODULAR

Use clear interfaces:


```
sensor/
preprocessing/
alignment/
ml/
navigation/
ekf/
constraints/
gnss/
mapping/
api/
simulation/
evaluation/
frontend/
```


Exact directory names may differ based on the existing repository.

The key requirement is separation of responsibilities.

45. FINAL END-TO-END PIPELINE

The working system should conceptually implement:





46. GNSS OUTAGE FLOW

When GNSS is healthy:

GNSS + IMU + AI

↓

EKF

↓

Navigation

When GNSS fails:

GNSS outage detection

↓

DR mode

↓

IMU + AI velocity

↓

EKF

↓

NHC + ZUPT

↓

Offline map matching

↓

Navigation continues

When GNSS returns:

GNSS reacquisition



Quality check



Innovation gating



Controlled EKF update



Smooth fusion



Normal GNSS-assisted navigation

47. FINAL VALIDATION

Do not say “complete” simply because the application launches.

Completion requires evidence that:

- backend starts
 - frontend connects
 - sensors are accepted
 - data is synchronized
 - preprocessing works
 - AI produces valid velocity
 - EKF remains stable
 - NHC works
 - ZUPT works
 - GNSS outage works
 - DR continues
 - map matching works
 - GNSS recovery works
 - UI reflects actual state
 - replay works
 - tests pass
 - performance is measured
 - offline behavior is verified
-

48. FINAL REPORT

After implementation, produce:

1. Initial audit

What existed.

2. Problems discovered

What was broken.

3. Root causes

Why it was broken.

4. Changes

What was modified.

5. Final architecture

Backend + ML + navigation + frontend.

6. Test results

Exact tests and results.

7. Benchmark results

Actual measurements.

8. Offline verification

Exactly how offline operation was tested.

9. Deployment status

Clearly state whether the current implementation is:

- development
- local demonstration
- smartphone deployment ready

Do not confuse these.

10. Remaining limitations

Everything not yet solved.

11. Requirement matrix

For every major requirement:

Requirement	Status	Evidence
Smartphone-only	PASS/FAIL/BLOCKED	Evidence
No external hardware	PASS/FAIL/BLOCKED	Evidence
Python backend	PASS/FAIL/BLOCKED	Evidence

Requirement	Status	Evidence
AI velocity	PASS/FAIL/BLOCKED	Evidence
15-state EKF	PASS/FAIL/BLOCKED	Evidence
NHC	PASS/FAIL/BLOCKED	Evidence
ZUPT	PASS/FAIL/BLOCKED	Evidence
GNSS outage	PASS/FAIL/BLOCKED	Evidence
GNSS recovery	PASS/FAIL/BLOCKED	Evidence
Offline map matching	PASS/FAIL/BLOCKED	Evidence
Offline runtime	PASS/FAIL/BLOCKED	Evidence
Performance	PASS/FAIL/BLOCKED	Measured evidence

49. FINAL COMMAND

Start immediately with:

PHASE 1 — REPOSITORY AUDIT

- Do not ask me to manually identify files that you can inspect yourself.
- Do not begin by rewriting the project.
- First understand the existing implementation.
- Then establish the baseline.
- Then create a prioritized repair plan.
- Then implement the fixes.
- Then integrate the complete pipeline.
- Then test.
- Then benchmark.
- Then fix failures.
- Then test again.
- Continue iterating until the system genuinely works or a technically demonstrated blocker prevents completion.
- The objective is NOT to make the repository look complete.
- The objective is to make **Navigators actually work as a coherent smartphone-only intelligent dead-reckoning navigation system.**
- No fake functionality.
- No fabricated metrics.
- No hidden dependencies.

No unnecessary hardware.

No cloud dependency.

No test cheating.

No premature completion claims.